Clustering and Similarity: Retrieving Documents

Emily Fox & Carlos Guestrin

Machine Learning Specialization

University of Washington

Retrieving documents of interest

Document retrieval

- Currently reading article you like

Document retrieval

- Currently reading article you like
- **Goal**: Want to find similar article

Document retrieval

Challenges

- How do we measure similarity?
- How do we search over articles?

是否有效率高的找法?
(暴力解需看過所有文章)

Word count representation for measuring similarity

Word count document representation

- Bag of words model
  - Ignore order of words
  - Count # of instances of each word in vocabulary

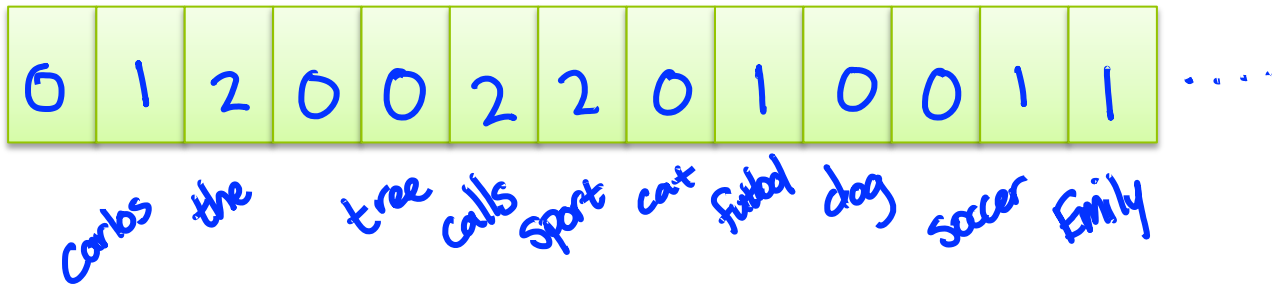

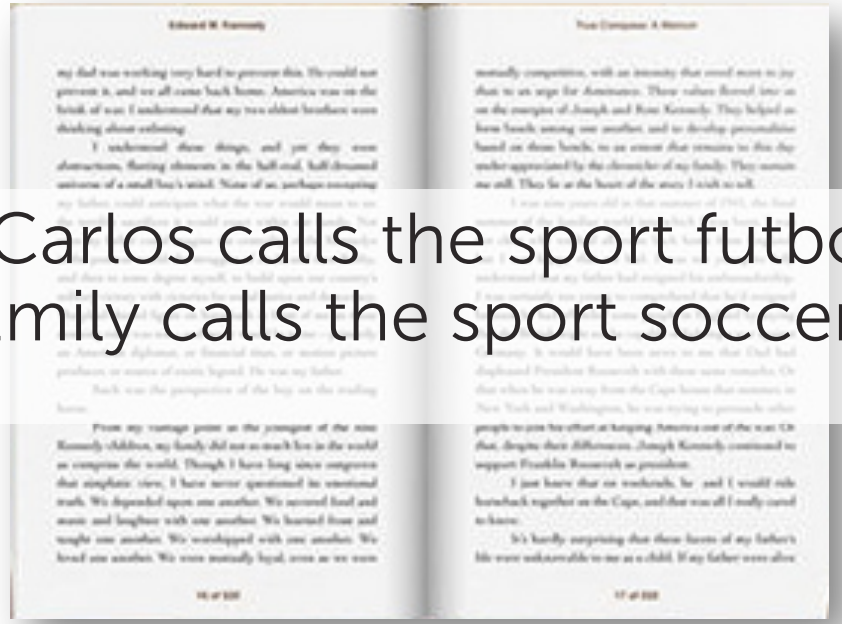

"Carlos calls the sport futbol. Emily calls the sport soccer."

Measuring similarity

相似度：向量內積

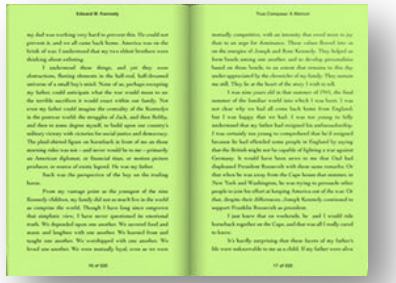

1	0	0	0	5	3	0	0	1	0	0	0	0

1*3
+
5*2
= **13**

3	0	0	0	2	0	0	1	0	1	0	0	0

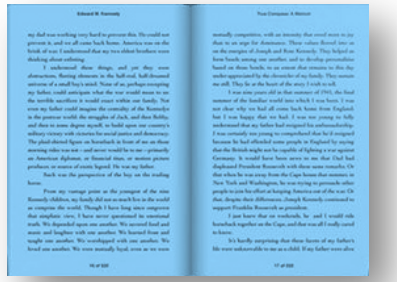

Measuring similarity

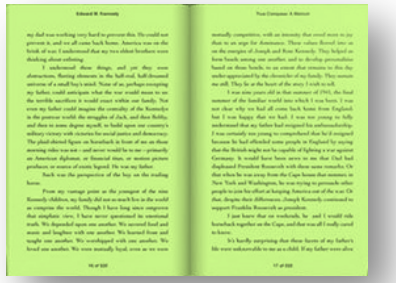

1	0	0	0	5	3	0	0	1	0	0	0	0

→ **0**

0	0	1	0	0	0	9	0	0	6	0	4	0

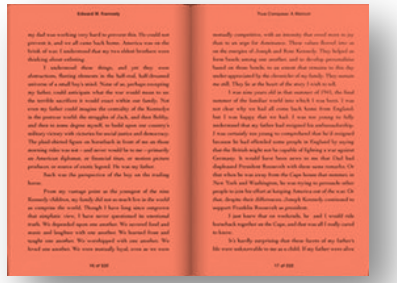

Issues with word counts – Doc length

用 word count 會受文章長度影響

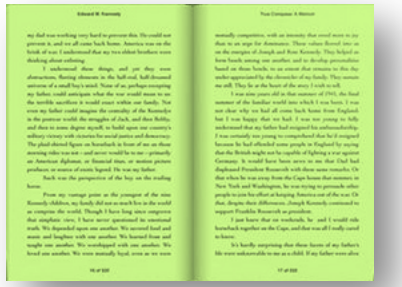

1	0	0	0	5	3	0	0	1	0	0	0	0

2	0	0	0	10	6	0	0	2	0	0	0	0

3	0	0	0	2	0	0	1	0	1	0	0	0

Similarity = 13

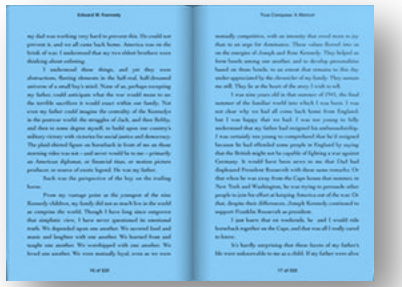

6	0	0	0	4	0	0	2	0	2	0	0	0

Similarity = 52

Solution = normalize

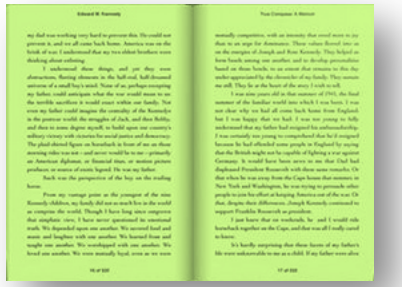

1	0	0	0	5	3	0	0	1	0	0	0	0

$$\sqrt{(1^2 + 5^2 + 3^2 + 1^2)}$$

1/6	0	0	0	5/6	3/6	0	0	1/6	0	0	0	0

Prioritizing important words with tf-idf

Issues with word counts – Rare words

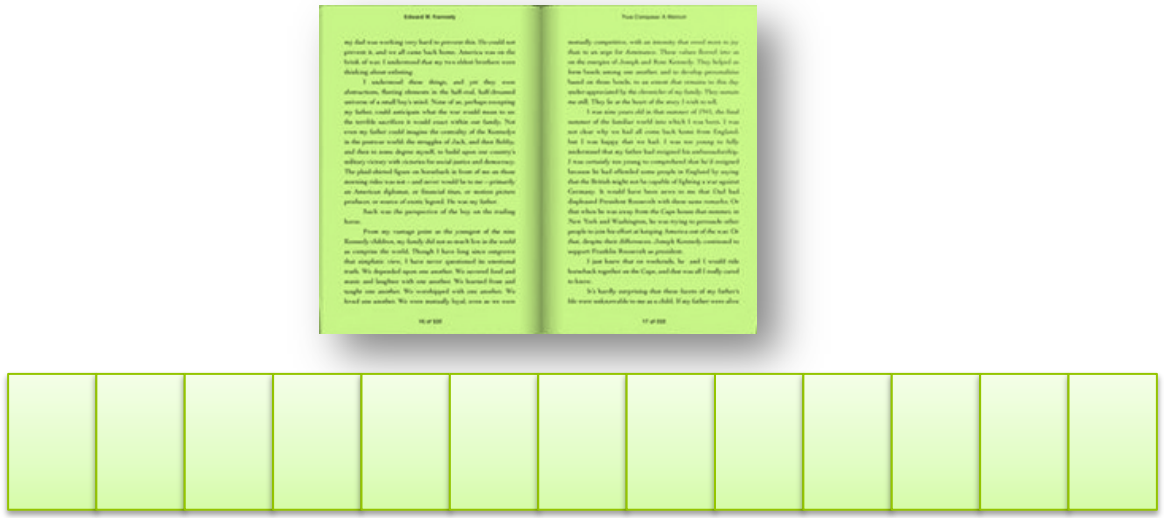

Common words in doc: “the”, “player”, “field”, “goal”

Dominate rare words like: “futbol”, “Messi”

Document frequency

- What characterizes a rare word?
  - Appears infrequently in the corpus

- Emphasize words appearing in few docs
  - Equivalently, discount word *w* based on
    # of docs containing *w* in corpus

Important words

- Do we want only rare words to dominate???

- What characterizes an important word?
  - Appears frequently in document (common locally)
  - Appears rarely in corpus (rare globally)

- Trade off between local frequency and global rarity

注意這是個假設
在單一 (少數) 文章裡重複出現且整體文章不常看到
=> 重要的字

TF-IDF document representation

- Term frequency – inverse document frequency (tf-idf)

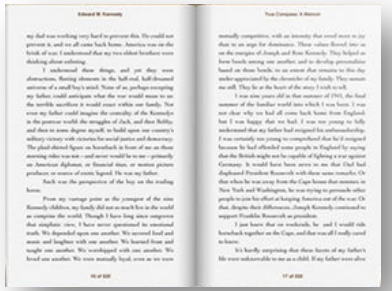

TF-IDF document representation

- Term frequency – inverse document frequency (tf-idf)
- Term frequency

- Same as word counts

其實有很多種表示法

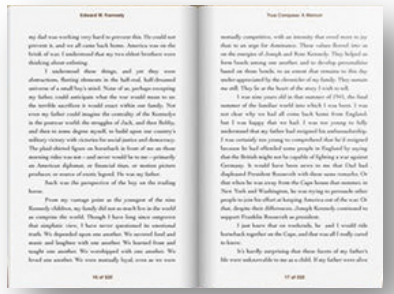

TF-IDF document representation

- Term frequency – inverse document frequency (tf-idf)
- Term frequency

- Inverse document frequency

$$\log \frac{\#\text{ docs}}{1 + \#\text{ docs using word}}$$

避免除以 0

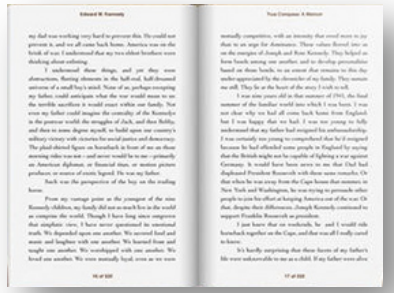

TF-IDF document representation

- Term frequency – inverse document frequency (tf-idf)
- Term frequency

- Inverse document frequency

$$\log \frac{\#\text{ docs}}{1 + \#\text{ docs using word}}$$

如果定義成這樣?

$$\frac{\#\text{doc}}{1 + \#\text{docs using word}} - 1$$

Why log?

$\log \frac{\text{large \#}}{1+\text{large \#}} \approx \log 1 = 0$

word in many docs

rare word

$\log \frac{\text{large \#}}{1+\text{small \#}} \rightarrow$ large #

TF-IDF document representation

- Term frequency – inverse document frequency (tf-idf)
- Term frequency
- Inverse document frequency

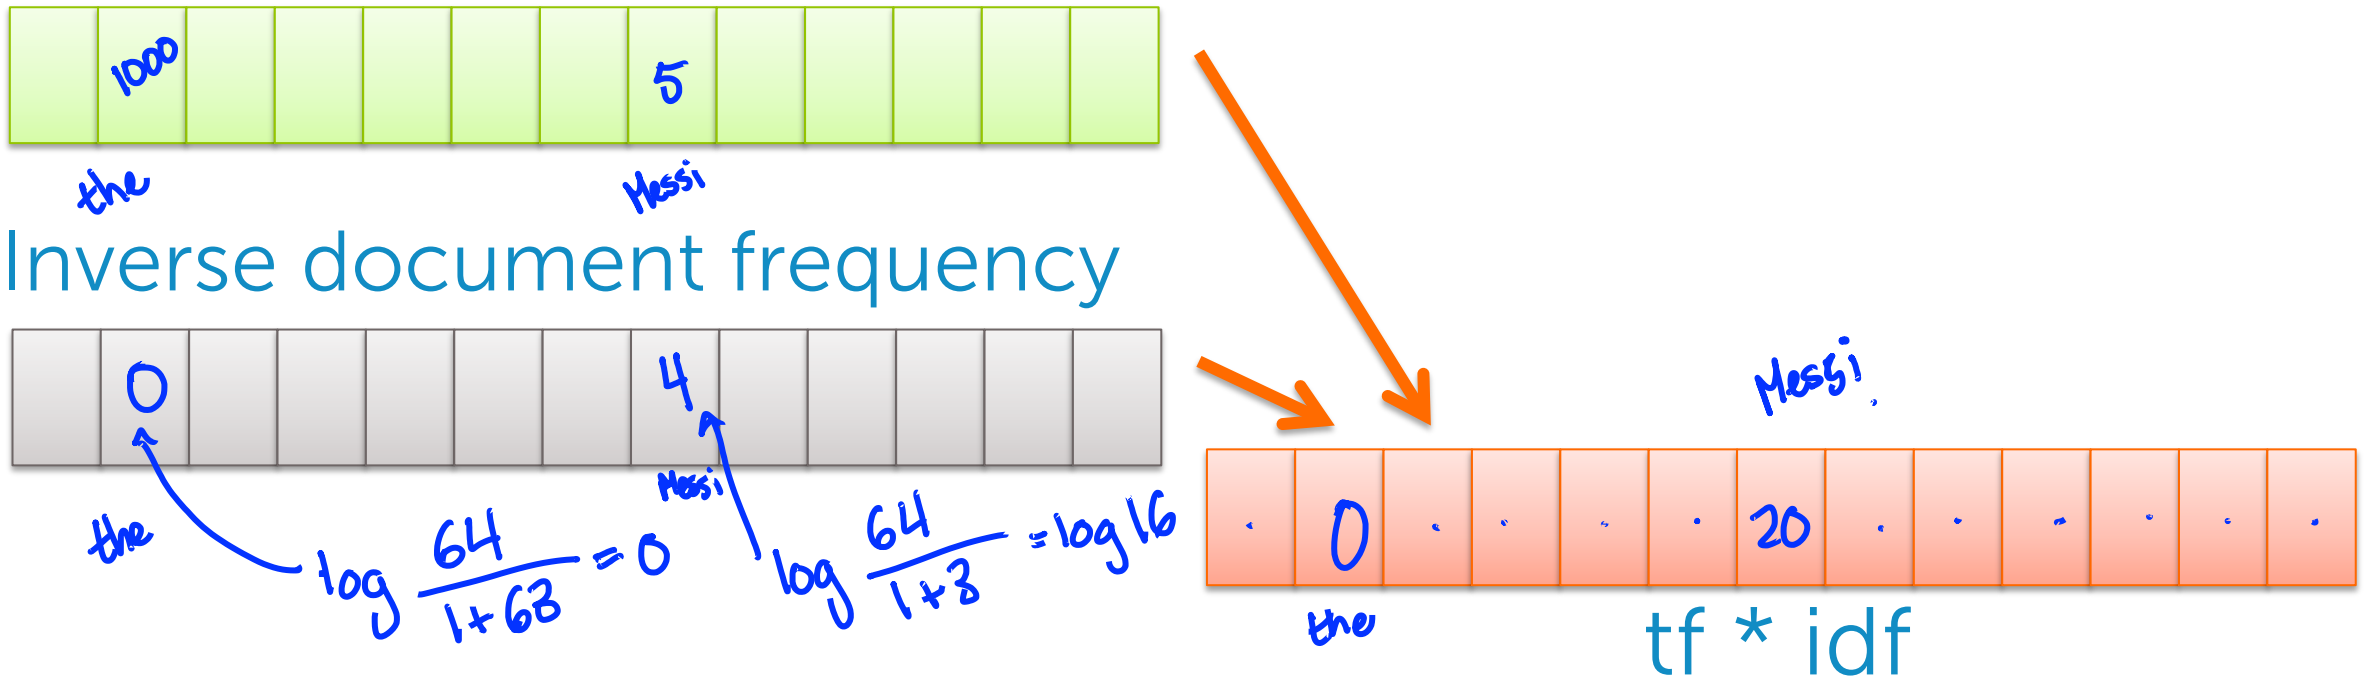

Retrieving similar documents

Nearest neighbor search

- Query article:

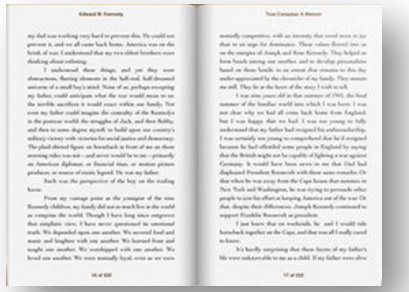

- Corpus:

- **Specify:** Distance metric

距離公式很可能影響結果

- **Output:** Set of most similar articles

1 – Nearest neighbor

找最相近的

假設?

- **Input:** Query article 
- **Output:** *Most* similar article

- Algorithm:
  - Search over each article  in corpus
    - Compute **s = similarity(** 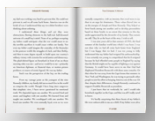 ,  **)**
    - If **s > Best_s**, record  =  and set **Best_s = s**
  - Return 

k – Nearest neighbor

找 k 個最相近的 假設?

- **Input:** Query article 
- **Output:** *List of k* similar articles

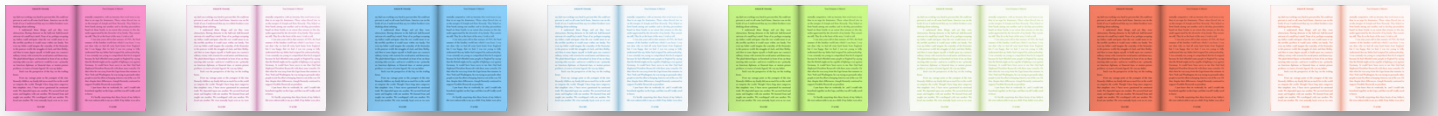

Clustering documents

參照 Andrew 的 machine learning， week 8 前半段

Structure documents by topic

- Discover groups (*clusters*) of related articles

SPORTS

WORLD NEWS

What if some of the labels are known?

- Training set of labeled docs

SPORTS

WORLD NEWS

ENTERTAINMENT

SCIENCE

Multiclass classification problem

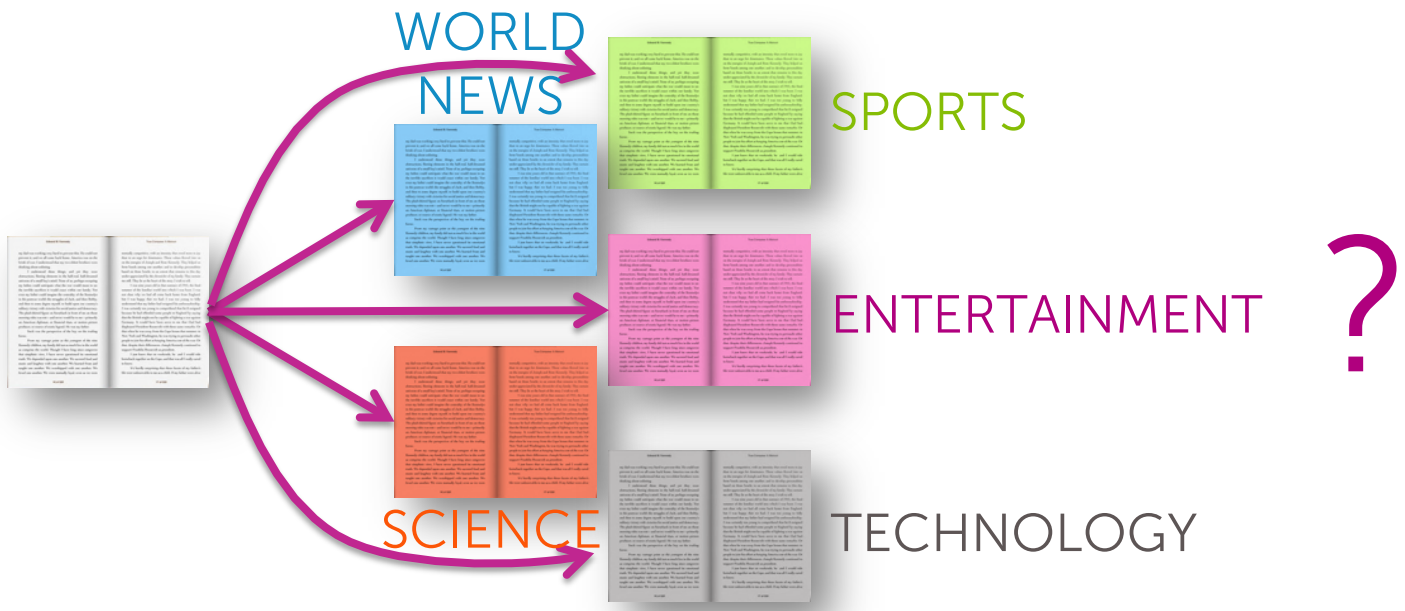

分群 (clustering) 的問題
在 training data label 已知的情形下可視為多類別的 supervised learning

Example of
supervised learning

Clustering

實際上沒有 label
需要找出結構

找出結構後，人為地
在事後給定 label

- No labels provided
- Want to uncover cluster structure

- **Input:** docs as vectors
- **Output:** cluster labels

An **unsupervised learning** task

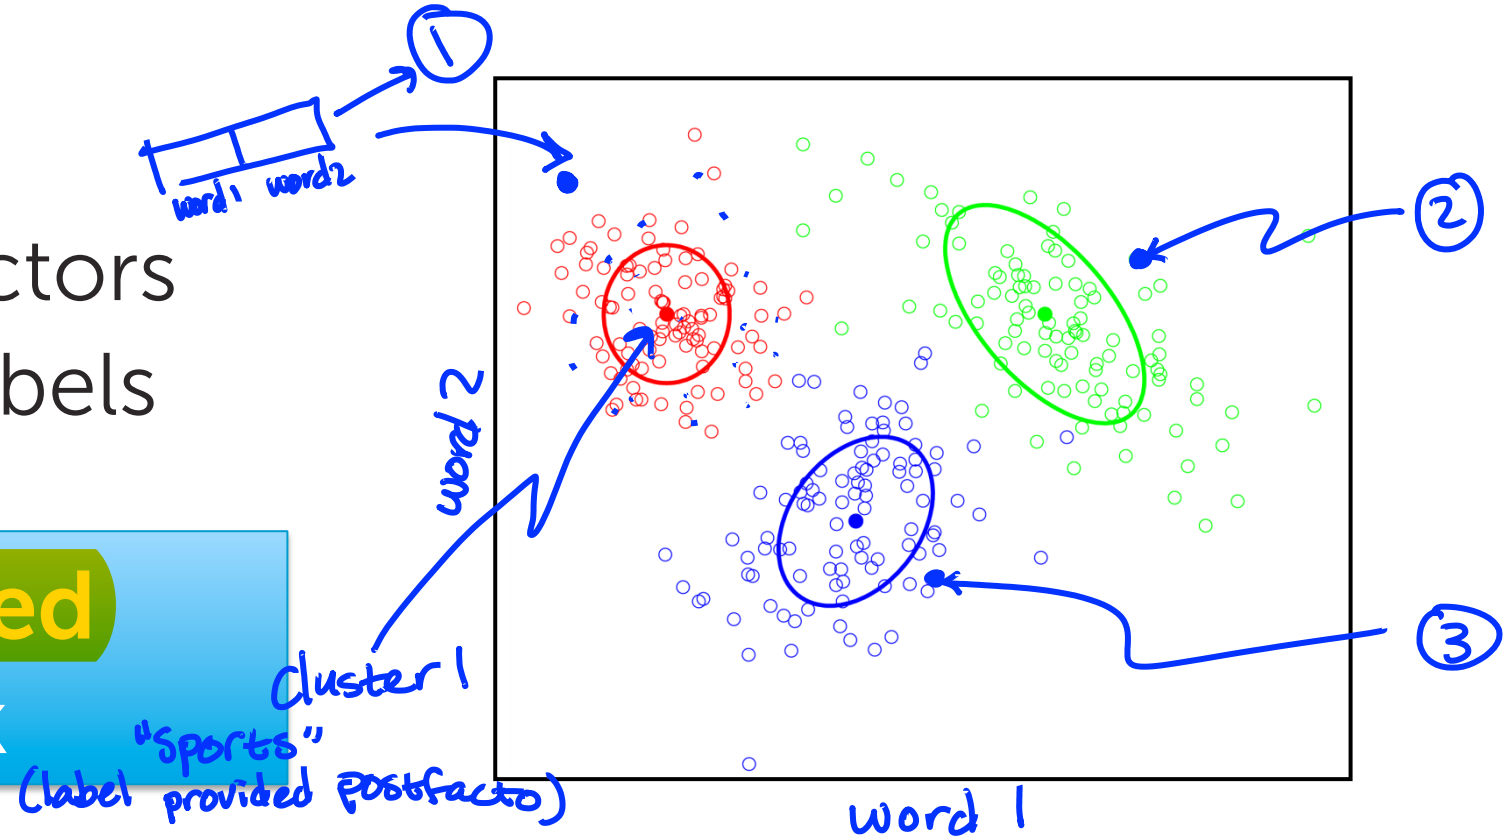

What defines a cluster?

- Cluster defined by center & shape/spread

別忽視形狀、分布

- Assign observation (doc) to cluster (topic label)
  - Score under cluster is higher than others
  - Often, just more similar to assigned cluster center than other cluster centers

只用跟中心點的距離，此點可能會被視為紅色或綠色

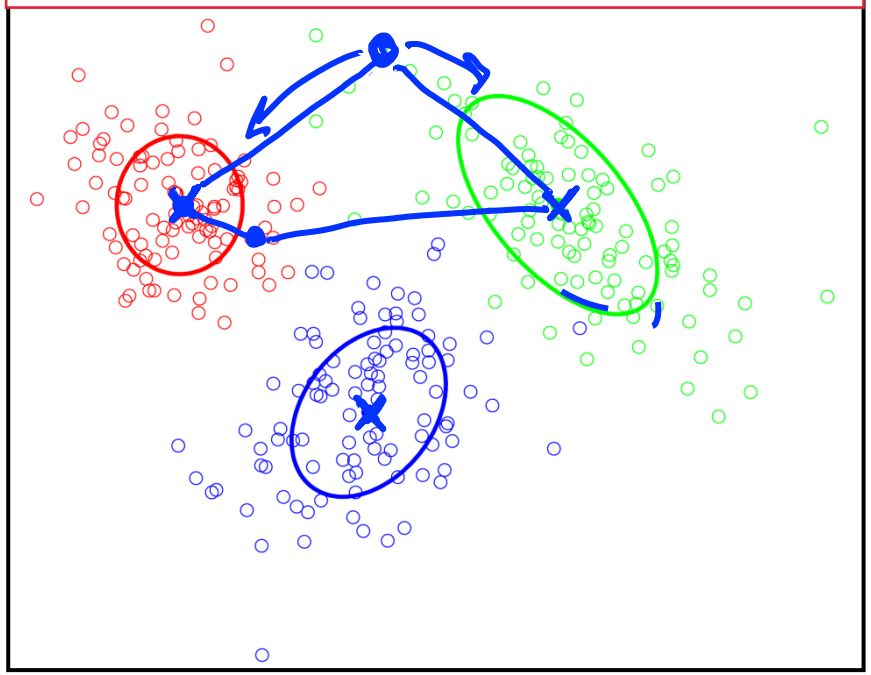

k-means

- Assume
  - Similarity metric = distance to cluster center (smaller better)

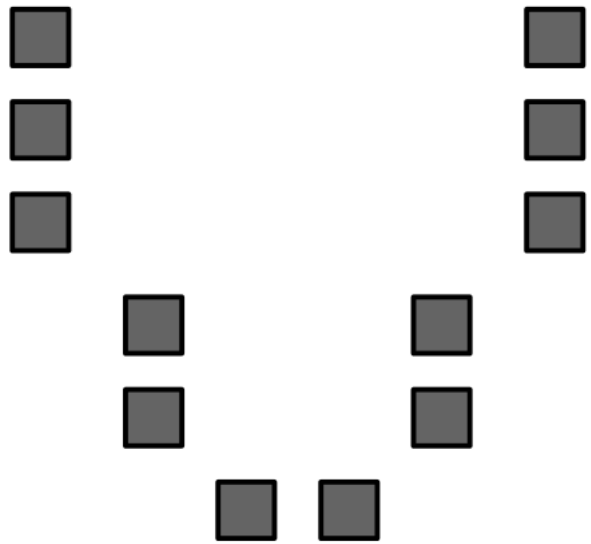

DATA to CLUSTER

k-means algorithm

0. Initialize cluster centers

此例子看起來是隨機生成中心點

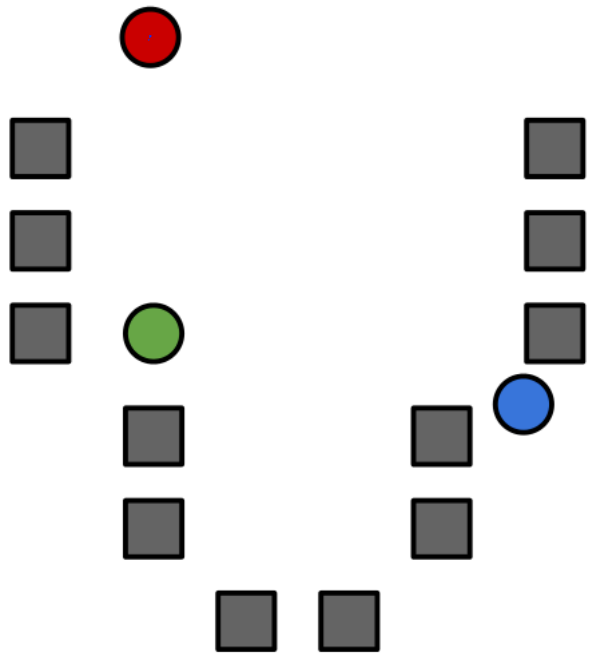

k-means algorithm

0. Initialize cluster centers
1. Assign observations to closest cluster center

還記得假設嗎?

這個示意圖有什麼問題?

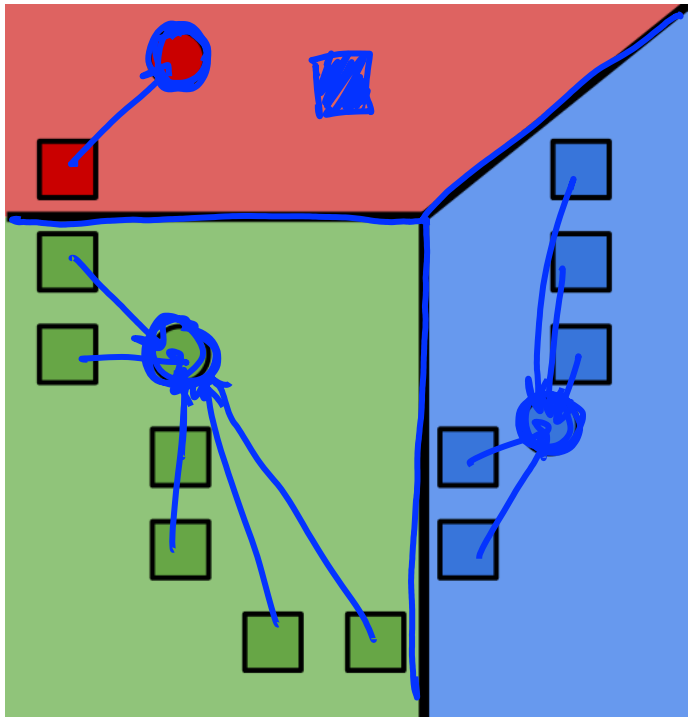

k-means algorithm

0. Initialize cluster centers
1. Assign observations to closest cluster center
2. Revise cluster centers as mean of assigned observations

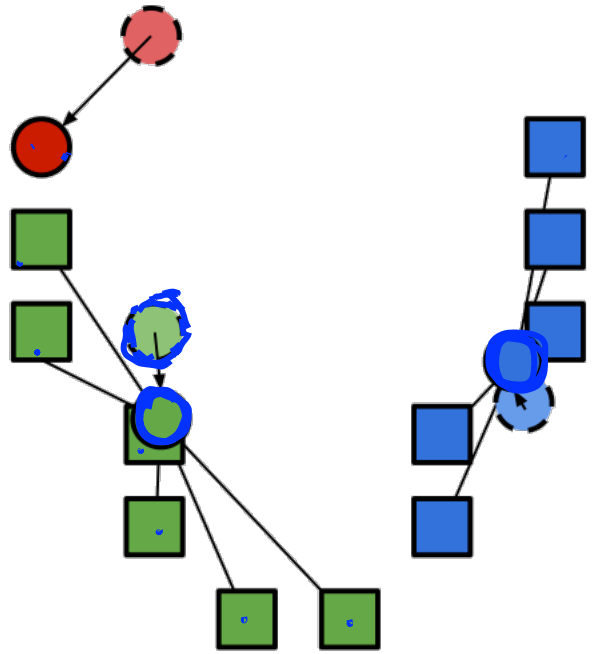

k-means algorithm

0. Initialize cluster centers
1. Assign observations to closest cluster center
2. Revise cluster centers as mean of assigned observations
3. Repeat 1.+2. until convergence

這次就忽視小瑕疵吧

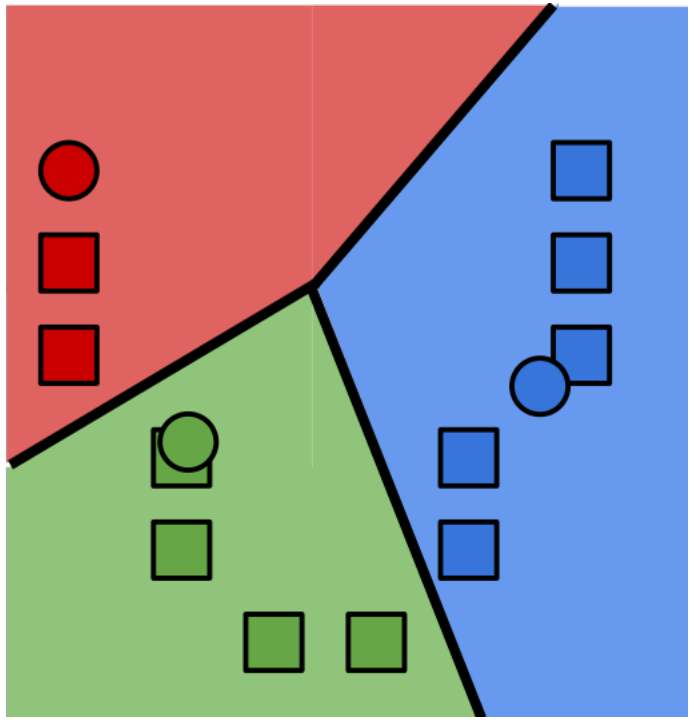

Other examples

後續無 note

Clustering images

- For search, group as:
  - Ocean
  - Pink flower
  - Dog
  - Sunset
  - Clouds
  - …

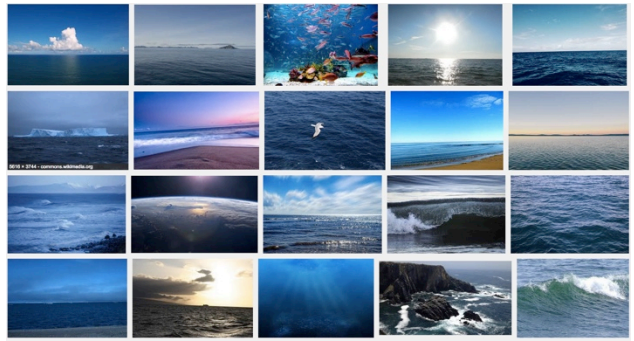

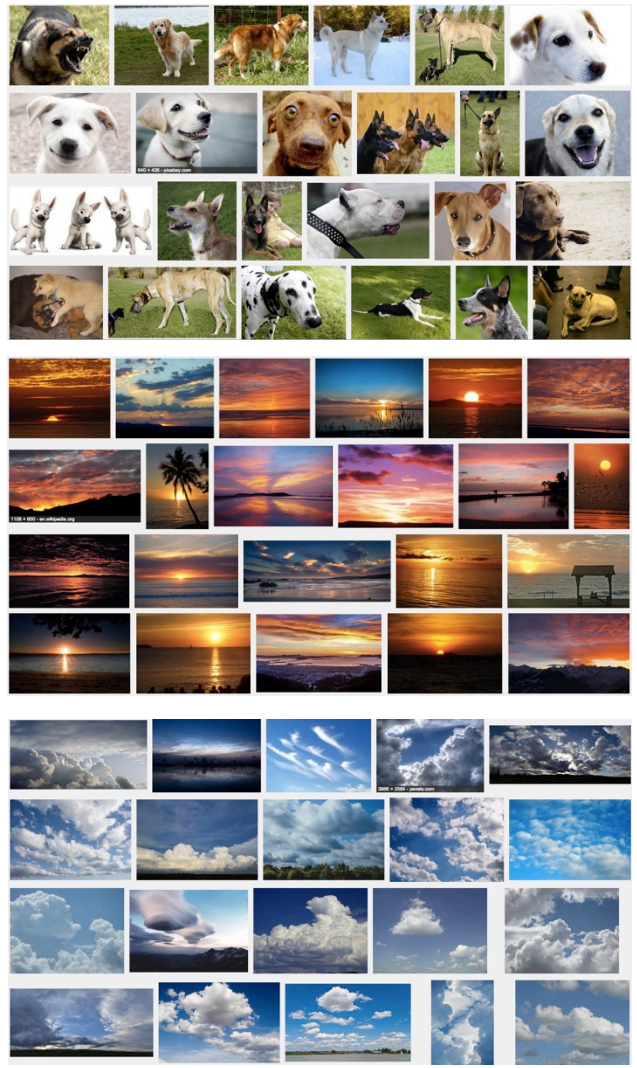

Grouping patients by medical condition

- Better characterize subpopulations and diseases

Example: Patients and seizures are diverse

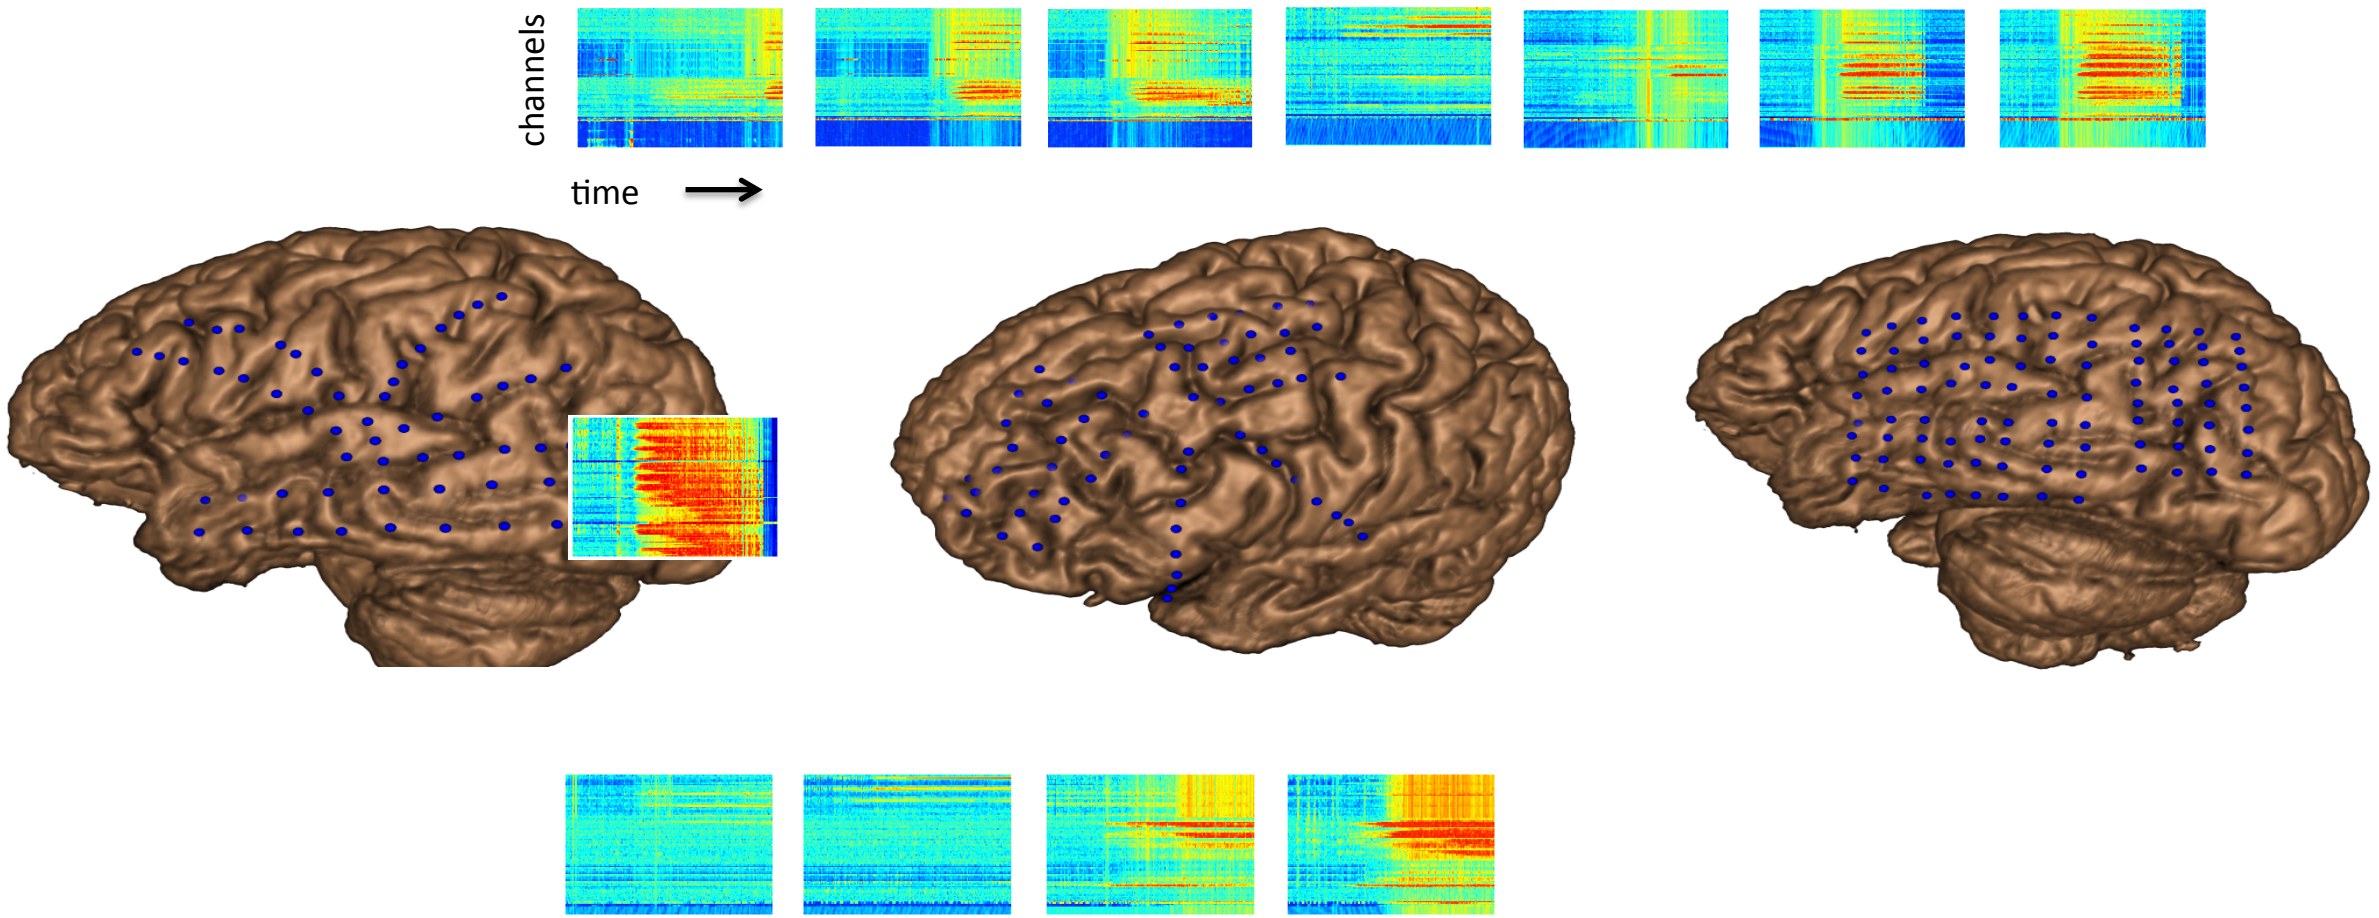

Cluster seizures by observed time courses

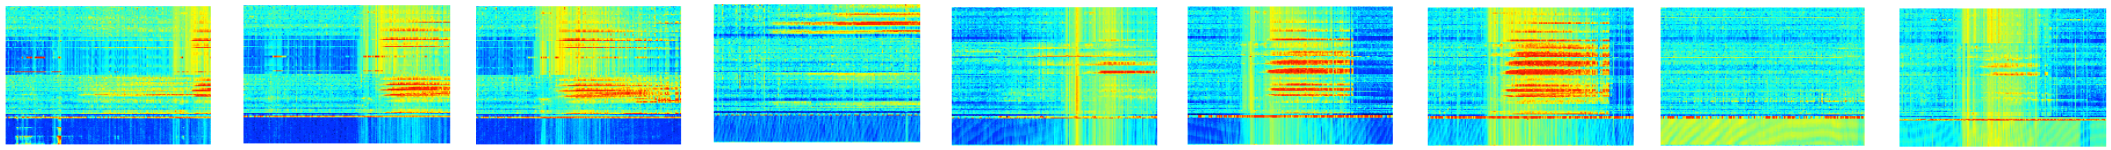

Products on Amazon

- Discover product categories from purchase histories

~~"furniture"~~

"baby"

- Or discovering groups of **users**

Structuring web search results

- Search terms can have multiple meanings
- Example: "cardinal"

- Use clustering to structure output

Discovering similar neighborhoods

- **Task 1:** Estimate price at a small regional level
- Challenge:
  - Only a few (or no!) sales in each region per month
- Solution:
  - Cluster regions with similar trends and share information within a cluster

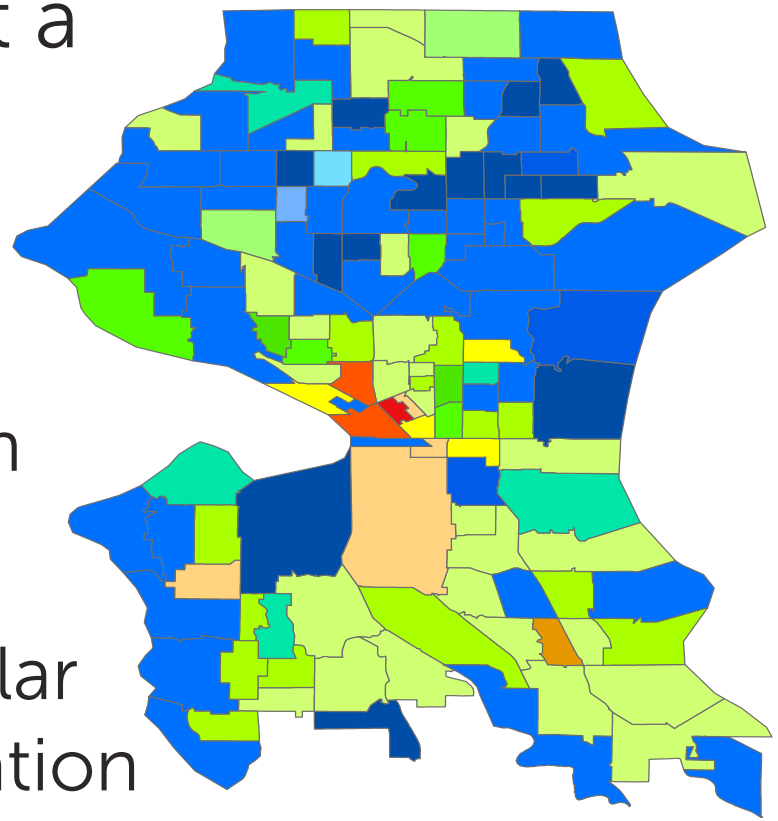

City of Seattle

Discovering similar neighborhoods

- **Task 2:** Forecast violent crimes to better task police
- Again, cluster regions and share information!
- Leads to improved predictions compared to examining each region independently

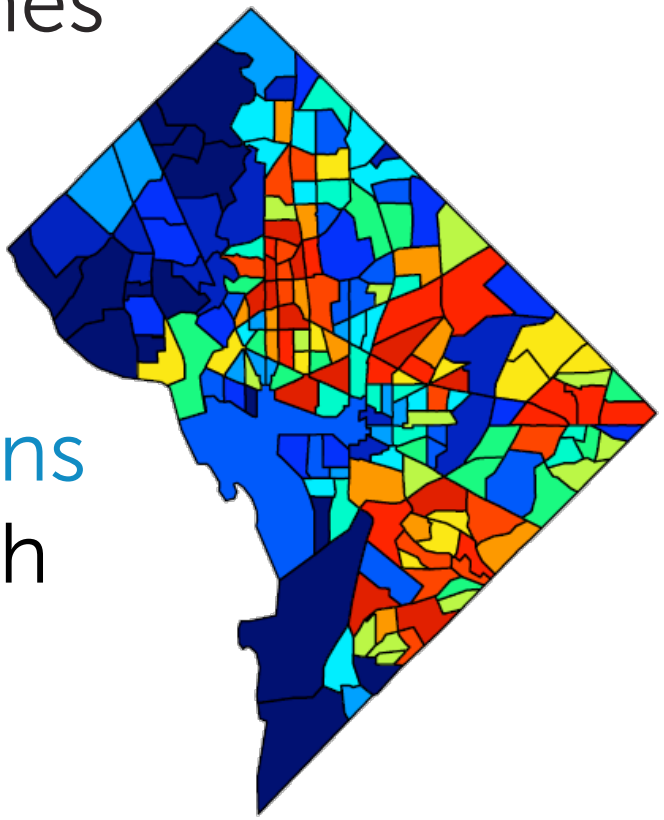

Washington, DC

Summary for clustering and similarity

What you can do now…

- Describe ways to represent a document (e.g., raw word counts, tf-idf,…)
- Measure the similarity between two documents
- Discuss issues related to using raw word counts
  - Normalize counts to adjust for document length
  - Emphasize important words using tf-idf
- Implement a nearest neighbor search for document retrieval
- Describe the input (unlabeled observations) and output (labels) of a clustering algorithm
- Determine whether a task is supervised or unsupervised
- Cluster documents using k-means (algorithmic details to come…)
- Describe other applications of clustering